

run_train_distracted_drivers-maxout8-size-261

August 4, 2016

The changes here are larger images but same maxout.

```
In [1]: from skimage import io, transform, exposure, color, util
import os, itertools, sys
from PIL import Image
%pylab inline
sys.setrecursionlimit(1000000)
```

Populating the interactive namespace from numpy and matplotlib

```
In [2]: # data_dir = "/home/dylan/IdeaProjects/distracted_drivers/train/"
data_dir = "/media/dylan/Science/Kaggle-Data/distracted_drivers/train/"
```

```
In [3]: input_volume_shape = (261, 261)
```

```
In [4]: def read_img_file_PIL(file_path, size=(32,32)):
    img = Image.open(file_path).convert('L')
    img.thumbnail(size, Image.NEAREST)
    data = np.array(img)
    shape = data.shape
    append_top = int(ceil(max(0, size[0] - shape[0])/2.0))
    append_bot = int(floor(max(0, size[0] - shape[0])/2.0))
    data = util.pad(data, ((append_top, append_bot),
                           (0,0)), mode='constant', constant_values=0)

    return data
```

```
In [5]: def read_img_file(file_path, rescale=0.01):
    img = io.imread(file_path)
    img = color.rgb2gray(img)
    return transform.rescale(img, rescale)
```

```
In [6]: def image_gen_from_dir(directory, batch_size, num_categories, size=input_volume_shape):
    result = {os.path.join(dp, f) : int(os.path.split(dp)[1]) for dp, dn, filenames in os.walk(
        directory)
        for f in filenames if os.path.splitext(f)[1] == '.jpg'}

    # infinite loop
    while True:
        image_files = []
        labels = []
        # randomly choose batch size samples in result
        for category in range(num_categories):
            file_samples = np.random.choice([k for k, v in result.iteritems() if v == category],
                                             size=batch_size, replace=False)
            for file_sample in file_samples:
                image_files.append(read_img_file_PIL(file_sample, size=size))
            labels.extend([category for v in itertools.repeat(category, batch_size)])
```

```

        # end category loop
    X = np.asarray(image_files, dtype=np.float32)
    # -1 to 1 range
    X = exposure.rescale_intensity(X, out_range=(-1,1))
    y = np.asarray(labels, dtype=np.int32)
    yield X, y

```

0.1 Another loader, augmentation time

We'll do 6 augmentations:

- 1.) Translation up to 10 pixels
- 2.) Rotation up to 15 degrees
- 3.) Zooming
- 4.) JPEG compression
- 5.) Sharpening
- 6.) Gamma correction

We won't do flips since the dataset only contains images from the passenger seat. Perhaps we can revisit this later.

```

In [7]: from skimage.transform import rotate, warp, AffineTransform
        from skimage import filters
        from scipy import ndimage, misc
        import StringIO

```

```

In [8]: def random_translate(img):
        shift_random = AffineTransform(translation=(randint(-10, 10), randint(-10, 10)))
        min_value = 0 if min(img.ravel()) > 0 else min(img.ravel())
        return np.float32(warp(img, shift_random, mode='constant', cval=min_value))

```

```

def random_rotate(img):
    min_value = 0 if min(img.ravel()) > 0 else min(img.ravel())
    return np.float32(rotate(img, randint(-15, 15), mode='constant', cval=min_value))

```

```

def random_zoom(img):
    min_value = 0 if min(img.ravel()) > 0 else min(img.ravel())
    scale_random = AffineTransform(scale=(uniform(0.9, 1.1), uniform(0.9, 1.1)))
    return np.float32(warp(img, scale_random, mode='constant', cval=min_value))

```

```

def random_compress(img):
    max_v = np.ceil(img.max())
    min_v = np.floor(img.min())
    nd_im = exposure.rescale_intensity(img, out_range=(0, 1)).squeeze()
    nd_im = np.ndarray.astype(nd_im * 255, np.uint8)
    # nd_im = np.ndarray.astype(img * 255, np.uint8)
    im = Image.fromarray(nd_im)
    buf = StringIO.StringIO()
    im.save(buf, "JPEG", quality=np.random.randint(95, 99))
    buf.seek(0)
    im2 = Image.open(buf)
    x1 = exposure.rescale_intensity(np.ndarray.astype(np.array(im2), np.float32), out_range=(min
    return x1

```

```

def random_sharpening(img):
    blurred_f = ndimage.gaussian_filter(img, 0.5)
    filter_blurred_f = ndimage.gaussian_filter(blurred_f, 1)
    alpha = uniform(0.9, 1.2)
    img = blurred_f + alpha * (blurred_f - filter_blurred_f)
    return exposure.rescale_intensity(img, out_range=(-1, 1))

def random_gamma_correction(img):
    max_v = np.ceil(img.max())
    min_v = np.floor(img.min())
    img = exposure.rescale_intensity(img, out_range=(0,1))
    img = exposure.adjust_gamma(img, uniform(0.2, 0.8))
    return exposure.rescale_intensity(img, out_range=(-1, 1))

In [9]: def random_aug(img):
    choice = np.random.randint(0,6)
    # choose from 4 different augmentations!
    if choice == 0:
        return random_translate(img)
    elif choice == 1:
        return random_rotate(img)
    elif choice == 2:
        return random_zoom(img)
    elif choice == 3:
        return random_compress(img)
    elif choice == 4:
        return random_sharpening(img)
    else:
        return random_gamma_correction(img)

In [10]: def random_aug_batch(X, aug_algorithm):
    for i in range(X.shape[0]):
        X[i] = aug_algorithm(X[i])
    return X

In [11]: def random_aug_gen(gen, aug_algorithm):
    for batchX, batchY in gen:
        yield random_aug_batch(batchX, aug_algorithm), batchY

```

1 Process Generator with cached elements

```

In [12]: def threaded_generator(generator, num_cached=50):
    import Queue
    queue = Queue.Queue(maxsize=num_cached)
    sentinel = object() # guaranteed unique reference

    # define producer (putting items into queue)
    def producer():
        for item in generator:
            queue.put(item)
        queue.put(sentinel)

    # start producer (in a background thread)
    import threading

```

```

thread = threading.Thread(target=producer)
thread.daemon = True
thread.start()

# run as consumer (read items from queue, in current thread)
item = queue.get()
while item is not sentinel:
    yield item
    queue.task_done()
    item = queue.get()

```

```

In [13]: from nolearn.lasagne import NeuralNet
        from lasagne.layers import DenseLayer, ReshapeLayer, Upscale2DLayer, Conv2DLayer, InputLayer,
        MaxPool2DLayer, get_all_params, batch_norm, BatchNormLayer, FeaturePoolLayer
        import numpy as np
        from lasagne.nonlinearities import softmax, leaky_rectify, theano
        from lasagne.updates import nesterov_momentum
        from nolearn.lasagne import NeuralNet, BatchIterator, PrintLayerInfo, objective
        from nolearn.lasagne import TrainSplit
        from common import EarlyStopping, EndTrainingFromEarlyStopping
        from lasagne.objectives import categorical_crossentropy, aggregate
        import cPickle as pickle
        from sklearn import metrics
        import time, logging, logging.config, logging.handlers
        from lasagne.init import Orthogonal
        from notebook_functions import load_best_weights

```

Couldn't import dot_parser, loading of dot files will not be possible.

Using gpu device 0: GeForce GTX 960 (CNMeM is disabled, CuDNN 4004)

```
def batch_norm(s): return s
```

In [14]: try:

```

from lasagne.layers.dnn import Conv2DDNNLayer, MaxPool2DDNNLayer
def conv_2_layer_stack(top, num_filters, pad=1):
    conv1 = batch_norm(Conv2DDNNLayer(top, num_filters, (3, 3),
        stride=1, pad=pad, nonlinearity=leaky_rectify, W=Orthogonal()))
    conv2 = batch_norm(Conv2DDNNLayer(conv1, num_filters, (3, 3),
        stride=1, pad=pad, nonlinearity=leaky_rectify, W=Orthogonal()))
    return MaxPool2DDNNLayer(conv2, (2, 2), 2)

def conv_3_layer_stack(top, num_filters, pad=1):
    conv1 = batch_norm(Conv2DDNNLayer(top, num_filters, (3, 3),
        stride=1, pad=pad, nonlinearity=leaky_rectify, W=Orthogonal()))
    conv2 = batch_norm(Conv2DDNNLayer(conv1, num_filters, (3, 3),
        stride=1, pad=pad, nonlinearity=leaky_rectify, W=Orthogonal()))
    conv3 = batch_norm(Conv2DDNNLayer(conv2, num_filters, (3, 3),
        stride=1, pad=pad, nonlinearity=leaky_rectify, W=Orthogonal()))
    return MaxPool2DDNNLayer(conv3, (2, 2), 2)

def conv_4_layer_stack(top, num_filters, pad=1):
    conv1 = batch_norm(Conv2DDNNLayer(top, num_filters, (3, 3),
        stride=1, pad=pad, nonlinearity=leaky_rectify, W=Orthogonal()))
    conv2 = batch_norm(Conv2DDNNLayer(conv1, num_filters, (3, 3),

```

```

        stride=1, pad=pad, nonlinearity=leaky_rectify, W=Orthogonal()))
conv3 = batch_norm(Conv2DDNNLayer(conv2, num_filters, (3, 3),
        stride=1, pad=pad, nonlinearity=leaky_rectify, W=Orthogonal()))
conv4 = batch_norm(Conv2DDNNLayer(conv3, num_filters, (3, 3),
        stride=1, pad=pad, nonlinearity=leaky_rectify, W=Orthogonal()))
return MaxPool2DDNNLayer(conv4, (2, 2), 2)

def conv_6_layer_stack(top, num_filters, pad=1):
conv1 = batch_norm(Conv2DDNNLayer(top, num_filters, (3, 3),
        stride=1, pad=pad, nonlinearity=leaky_rectify, W=Orthogonal()))
conv2 = batch_norm(Conv2DDNNLayer(conv1, num_filters, (3, 3),
        stride=1, pad=pad, nonlinearity=leaky_rectify, W=Orthogonal()))
conv3 = batch_norm(Conv2DDNNLayer(conv2, num_filters, (3, 3),
        stride=1, pad=pad, nonlinearity=leaky_rectify, W=Orthogonal()))
conv4 = batch_norm(Conv2DDNNLayer(conv3, num_filters, (3, 3),
        stride=1, pad=pad, nonlinearity=leaky_rectify, W=Orthogonal()))
conv5 = batch_norm(Conv2DDNNLayer(conv4, num_filters, (3, 3),
        stride=1, pad=pad, nonlinearity=leaky_rectify, W=Orthogonal()))
conv6 = batch_norm(Conv2DDNNLayer(conv5, num_filters, (3, 3),
        stride=1, pad=pad, nonlinearity=leaky_rectify, W=Orthogonal()))
return MaxPool2DLayer(conv6, (2, 2), 2)

except ImportError:
def conv_2_layer_stack(top, num_filters, pad=1):
conv1 = batch_norm(Conv2DLayer(
    top, num_filters, (3, 3), stride=1, pad=pad, nonlinearity=leaky_rectify))
conv2 = batch_norm(Conv2DLayer(
    conv1, num_filters, (3, 3), stride=1, pad=pad, nonlinearity=leaky_rectify))
return MaxPool2DLayer(conv2, (2, 2), 2)

def conv_3_layer_stack(top, num_filters, pad=1):
conv1 = batch_norm(Conv2DLayer(
    top, num_filters, (3, 3), stride=1, pad=pad, nonlinearity=leaky_rectify))
conv2 = batch_norm(Conv2DLayer(
    conv1, num_filters, (3, 3), stride=1, pad=pad, nonlinearity=leaky_rectify))
conv3 = batch_norm(Conv2DLayer(conv2
    , num_filters, (3, 3), stride=1, pad=pad, nonlinearity=leaky_rectify))
return MaxPool2DLayer(conv3, (2, 2), 2)

def conv_4_layer_stack(top, num_filters, pad=1):
conv1 = batch_norm(Conv2DLayer(
    top, num_filters, (3, 3), stride=1, pad=pad, nonlinearity=leaky_rectify))
conv2 = batch_norm(Conv2DLayer(
    conv1, num_filters, (3, 3), stride=1, pad=pad, nonlinearity=leaky_rectify))
conv3 = batch_norm(Conv2DLayer(
    conv2, num_filters, (3, 3), stride=1, pad=pad, nonlinearity=leaky_rectify))
conv4 = batch_norm(Conv2DLayer(
    conv3, num_filters, (3, 3), stride=1, pad=pad, nonlinearity=leaky_rectify))
return MaxPool2DLayer(conv4, (2, 2), 2)

def conv_6_layer_stack(top, num_filters, pad=1):
conv1 = batch_norm(Conv2DLayer(
    top, num_filters, (3, 3), stride=1, pad=pad, nonlinearity=leaky_rectify))
conv2 = batch_norm(Conv2DLayer(

```

```

        conv1, num_filters, (3, 3), stride=1, pad=pad, nonlinearity=leaky_rectify))
conv3 = batch_norm(Conv2DLayer(
    conv2, num_filters, (3, 3), stride=1, pad=pad, nonlinearity=leaky_rectify))
conv4 = batch_norm(Conv2DLayer(
    conv3, num_filters, (3, 3), stride=1, pad=pad, nonlinearity=leaky_rectify))
conv5 = batch_norm(Conv2DLayer(
    conv4, num_filters, (3, 3), stride=1, pad=pad, nonlinearity=leaky_rectify))
conv6 = batch_norm(Conv2DLayer(
    conv5, num_filters, (3, 3), stride=1, pad=pad, nonlinearity=leaky_rectify))
return MaxPool2DLayer(conv6, (2, 2), 2)

In [15]: k = 8
input_layer = InputLayer((None, 1, input_volume_shape[0], input_volume_shape[1]))
conv1 = batch_norm(Conv2DDNNLayer(input_layer, 32, (7, 7),
    stride=2, pad=0, nonlinearity=leaky_rectify, W=Orthogonal()))
maxpool1 = MaxPool2DLayer(conv1, (2, 2), 2)

conv_stack_1 = conv_2_layer_stack(maxpool1, 64)
# dropout1 = DropoutLayer(conv_stack_1, p=0.1)

conv_stack_2 = conv_2_layer_stack(conv_stack_1, 128)
# dropout2 = DropoutLayer(conv_stack_2, p=0.2)

conv_stack_3 = conv_2_layer_stack(conv_stack_2, 256)
# dropout3 = DropoutLayer(conv_stack_3, p=0.3)

# conv_stack_4 = conv_2_layer_stack(conv_stack_3, 512, pad=0)
# dropout4 = DropoutLayer(conv_stack_4, p=0.4)

# conv_stack_5 = conv_2_layer_stack(conv_stack_4, 1024, pad=0)
# dropout17 = DropoutLayer(conv_stack_5, p=0.5)

conv4 = batch_norm(Conv2DDNNLayer(conv_stack_3, 512, (3, 3),
    stride=1, pad=0, nonlinearity=leaky_rectify, W=Orthogonal()))
conv5 = batch_norm(Conv2DDNNLayer(conv4, 512, (3, 3),
    stride=1, pad=1, nonlinearity=leaky_rectify, W=Orthogonal()))
conv6 = batch_norm(Conv2DDNNLayer(conv5, 512, (3, 3),
    stride=1, pad=0, nonlinearity=leaky_rectify, W=Orthogonal()))
conv7 = batch_norm(Conv2DDNNLayer(conv6, 512, (3, 3),
    stride=1, pad=1, nonlinearity=leaky_rectify, W=Orthogonal()))
maxpool2 = MaxPool2DLayer(conv7, (2, 2), 2)

dense18 = DenseLayer(maxpool2, 2048, nonlinearity=None)
norm1 = BatchNormLayer(dense18)
maxout1 = FeaturePoolLayer(norm1, k)
dropout19 = DropoutLayer(maxout1, p=0.5)

dense20 = DenseLayer(dropout19, 2048, nonlinearity=None)
norm2 = BatchNormLayer(dense20)
maxout2 = FeaturePoolLayer(norm2, k)

softmax21 = DenseLayer(maxout2, 10, nonlinearity=softmax)

```

1.1 Quality of Life Functions

```
In [16]: if not os.path.exists("logs"):
        os.mkdir("logs")
        logging.config.fileConfig("logging-training.conf")

def regularization_objective(layers, lambda1=0., lambda2=0., *args, **kwargs):
    # default loss
    losses = objective(layers, *args, **kwargs)
    # get layer weights except for the biases
    weights = get_all_params(layers[-1], regularizable=True)
    regularization_term = 0.0
    # sum of abs weights for L1 regularization
    if lambda1 != 0.0:
        sum_abs_weights = sum([abs(w).sum() for w in weights])
        regularization_term += (lambda1 * sum_abs_weights)
    # sum of squares (sum(theta^2))
    if lambda2 != 0.0:
        sum_squared_weights = (1 / 2.0) * sum([(w ** 2).sum() for w in weights])
        regularization_term += (lambda2 * sum_squared_weights)
    # add weights to regular loss
    losses += regularization_term
    return losses

def eval_regularization(net):
    if net.objective_lambda1 == 0 and net.objective_lambda2 == 0:
        return 0
    # check the loss if the regularization term is not overpowering the loss
    weights = get_all_params(net.layers[-1], regularizable=True)
    # sum of abs weights for L1 regularization
    sum_abs_weights = sum([abs(w).sum() for w in weights])
    # sum of squares (sum(theta^2))
    sum_squared_weights = (1 / 2.0) * sum([(w ** 2).sum() for w in weights])
    # add weights to regular loss
    regularization_term = (net.objective_lambda1 * sum_abs_weights) \
        + (net.objective_lambda2 * sum_squared_weights)
    return regularization_term

def print_regularization_term(net):
    if net.objective_lambda1 > 0.0 or net.objective_lambda2 > 0.0:
        regularization_term = eval_regularization(net)
        print "Regularization term: {}".format(regularization_term.eval())

def validation_set_loss(_net, _X, _y):
    """We need this to track the validation loss"""
    _yb = _net.predict_proba(_X)
    _y_pred = np.argmax(_yb, axis=1)
    _acc = metrics.accuracy_score(_y, _y_pred)
    loss = aggregate(categorical_crossentropy(_yb, _y))
    loss += eval_regularization(_net)
    return loss, _acc
```

```

def store_model(model_file_name, net):
    directory_name = os.path.dirname(model_file_name)
    model_file_name = os.path.basename(model_file_name)
    if not os.path.exists(directory_name):
        os.makedirs(directory_name)
    # write model
    output_model_file_name = os.path.join(directory_name, model_file_name)
    start_write_time = time.time()
    if os.path.isfile(output_model_file_name):
        os.remove(output_model_file_name)
    with open(output_model_file_name, 'wb') as experiment_model:
        pickle.dump(net, experiment_model)
    total_write_time = time.time() - start_write_time
    m, s = divmod(total_write_time, 60)
    h, m = divmod(m, 60)
    logging.log(logging.INFO, "Duration of saving to disk: %0d:%02d:%02d", h, m, s)

def write_validation_loss_and_store_best(validation_file_name, best_weights_file_name,
                                         net, X_val, y_val, best_vloss, best_acc):
    # write validation loss
    start_validate_time = time.time()
    vLoss, vAcc = validation_set_loss(net, X_val, y_val)
    loss = vLoss.eval()
    current_epoch = net.train_history_[-1]['epoch']
    with open(validation_file_name, 'a') as validation_file:
        validation_file.write("{} , {}, {} \n".format(current_epoch, loss, vAcc))

    total_validate_time = time.time() - start_validate_time
    m, s = divmod(total_validate_time, 60)
    h, m = divmod(m, 60)
    logging.log(logging.INFO, "Duration of validation: %0d:%02d:%02d", h, m, s)

    # store best weights here
    if loss < best_vloss:
        start_bw_time = time.time()
        best_vloss = loss
        best_acc = vAcc
        with open(best_weights_file_name, 'wb') as best_model_file:
            pickle.dump(net.get_all_params_values(), best_model_file, -1)

    return best_vloss, best_acc

class AdjustVariableWithStepSize(object):
    """This class adjusts any variable during training"""

    def __init__(self, name, start=0.03, steps=3, after_epochs=2000):
        self.name = name
        self.start = start
        self.steps=steps
        self.after_epochs=after_epochs
        self.ls = []

```



```

def __call__(self, nn, train_history):
    if not self.ls:
        for i in range(self.steps):
            self.ls.extend(np.repeat(self.start/(np.power(10,i)), self.after_epochs))

    try:
        epoch = train_history[-1]['epoch']
        new_value = np.float32(self.ls[epoch - 1])
        getattr(nn, self.name).set_value(new_value)
    except IndexError:
        pass

```

1.2 CNN

```

In [17]: lambda1 = 0.0
         lambda2 = 5e-3

```

```

net = NeuralNet(
    layers=softmax21,
    max_epochs=1,
    update=nesterov_momentum,
    update_learning_rate=theano.shared(np.float32(0.001)),
    update_momentum = 0.99,
    # update=adam,
    on_epoch_finished=[
        EarlyStopping(patience=1000),
        AdjustVariableWithStepSize('update_learning_rate', start=0.001, steps=2, after_epochs=1),
    ],
    on_training_finished=[
        EndTrainingFromEarlyStopping()
    ],
    objective=regularization_objective,
    objective_lambda2=lambda2,
    objective_lambda1=lambda1,
    batch_iterator_train=BatchIterator(batch_size=100),
    train_split=TrainSplit(
        eval_size=0.25),
    # train_split=TrainSplit(eval_size=0.0),
    verbose=3,
)

```

```

In [18]: p = PrintLayerInfo()
         net.initialize()
         # p(net)

```

1.2.1 load cnn instead

```

In [19]: dir_name = 'net.vgg.large.12.5e3.orthog-norm-maxout8-lr.2.steps-size-261'
         validation_file_name = "{}/vloss-{}.txt".format(dir_name, dir_name)
         model_file_name = "{}/{}.pickle".format(dir_name, dir_name)
         best_weights_file_name = "{}/bw-{}.weights".format(dir_name, dir_name)
         if os.path.exists(dir_name):
             print "Model exists. Loading {}".format(dir_name)
             with open(model_file_name, 'rb') as reader:
                 net = pickle.load(reader)

```

```

else:
    print "Training model from the beginning {}".format(dir_name)

Training model from the beginning net.vgg.large.l2.5e3.orthog-norm-maxout8-lr.2.steps-size-261

load_best_weights(best_weights_file_name, net)
from nolearn.lasagne.visualize import plot_loss plt.figure( figsize=(15,9)) plt.ylim([0.1,0.5])
plt.plot([v['valid_loss'] for v in net.train_history_])

```

2 just this time.

```
net.on_epoch_finished.pop(1) print net.on_epoch_finished net.update_learning_rate=0.001
```

2.1 Define validation set

```

In [ ]: val_dir = "/media/dylan/Science/Kaggle-Data/distracted_drivers/val/"
        X_val, y_val = image_gen_from_dir(val_dir, 40, 10, size=input_volume_shape).next()
        X_val = X_val.reshape(-1, 1, input_volume_shape[0], input_volume_shape[1])

In [ ]: image_gen = image_gen_from_dir(data_dir, 10, 10, size=input_volume_shape)
        gen = random_aug_gen(image_gen, random_aug)
        threaded_gen = threaded_generator(gen, num_cached=100)

ops_every = 500
best_acc = 0.0
best_vloss = np.inf

start_time = time.time()
try:
    for step, (inputs, targets) in enumerate(threaded_gen):
        shape = inputs.shape
        net.fit(inputs.reshape(shape[0],1, shape[1], shape[2]), targets)
        logging.log(logging.INFO, "Epoch: {}, Training error: {}".format(
            net.train_history_-1["epoch"], net.train_history_-1["train_loss"]))
        if (step + 1) % ops_every == 0:
            print_regularization_term(net)
            store_model(model_file_name, net)
            # center validation
            best_vloss, best_acc = write_validation_loss_and_store_best(
                validation_file_name, best_weights_file_name, net, X_val, y_val, best_vloss, be

except StopIteration:
    # terminate if already early stopping
    with open(model_file_name, 'wb') as writer:
        pickle.dump(net, writer)
    total_time = time.time() - start_time
    print("Training successful by early stopping. Elapsed: {}".format(total_time))

# Neural Network with 14159722 learnable parameters

## Layer information

```

name	size	total	cap.Y	cap.X	cov.Y	cov.X	filter Y	filter X	f
InputLayer	1x261x261	68121	100.00	100.00	100.00	100.00	261	261	

Conv2DDNNLayer	32x128x128	524288	100.00	100.00	2.68	2.68	7	7
BatchNormLayer	32x128x128	524288	100.00	100.00	100.00	100.00	261	261
NonlinearityLayer	32x128x128	524288	100.00	100.00	100.00	100.00	261	261
MaxPool2DLayer	32x64x64	131072	100.00	100.00	100.00	100.00	261	261
Conv2DDNNLayer	64x64x64	262144	100.00	100.00	100.00	100.00	261	261
BatchNormLayer	64x64x64	262144	100.00	100.00	100.00	100.00	261	261
NonlinearityLayer	64x64x64	262144	100.00	100.00	100.00	100.00	261	261
Conv2DDNNLayer	64x64x64	262144	100.00	100.00	100.00	100.00	261	261
BatchNormLayer	64x64x64	262144	100.00	100.00	100.00	100.00	261	261
NonlinearityLayer	64x64x64	262144	100.00	100.00	100.00	100.00	261	261
MaxPool2DDNNLayer	64x32x32	65536	100.00	100.00	100.00	100.00	261	261
Conv2DDNNLayer	128x32x32	131072	100.00	100.00	100.00	100.00	261	261
BatchNormLayer	128x32x32	131072	100.00	100.00	100.00	100.00	261	261
NonlinearityLayer	128x32x32	131072	100.00	100.00	100.00	100.00	261	261
Conv2DDNNLayer	128x32x32	131072	100.00	100.00	100.00	100.00	261	261
BatchNormLayer	128x32x32	131072	100.00	100.00	100.00	100.00	261	261
NonlinearityLayer	128x32x32	131072	100.00	100.00	100.00	100.00	261	261
MaxPool2DDNNLayer	128x16x16	32768	100.00	100.00	100.00	100.00	261	261
Conv2DDNNLayer	256x16x16	65536	100.00	100.00	100.00	100.00	261	261
BatchNormLayer	256x16x16	65536	100.00	100.00	100.00	100.00	261	261
NonlinearityLayer	256x16x16	65536	100.00	100.00	100.00	100.00	261	261
Conv2DDNNLayer	256x16x16	65536	100.00	100.00	100.00	100.00	261	261
BatchNormLayer	256x16x16	65536	100.00	100.00	100.00	100.00	261	261
NonlinearityLayer	256x16x16	65536	100.00	100.00	100.00	100.00	261	261
MaxPool2DDNNLayer	256x8x8	16384	100.00	100.00	100.00	100.00	261	261
Conv2DDNNLayer	512x6x6	18432	100.00	100.00	100.00	100.00	261	261
BatchNormLayer	512x6x6	18432	100.00	100.00	100.00	100.00	261	261
NonlinearityLayer	512x6x6	18432	100.00	100.00	100.00	100.00	261	261
Conv2DDNNLayer	512x6x6	18432	100.00	100.00	100.00	100.00	261	261
BatchNormLayer	512x6x6	18432	100.00	100.00	100.00	100.00	261	261
NonlinearityLayer	512x6x6	18432	100.00	100.00	100.00	100.00	261	261
Conv2DDNNLayer	512x4x4	8192	100.00	100.00	100.00	100.00	261	261
BatchNormLayer	512x4x4	8192	100.00	100.00	100.00	100.00	261	261
NonlinearityLayer	512x4x4	8192	100.00	100.00	100.00	100.00	261	261
Conv2DDNNLayer	512x4x4	8192	100.00	100.00	100.00	100.00	261	261
BatchNormLayer	512x4x4	8192	100.00	100.00	100.00	100.00	261	261
NonlinearityLayer	512x4x4	8192	100.00	100.00	100.00	100.00	261	261
MaxPool2DLayer	512x2x2	2048	100.00	100.00	100.00	100.00	261	261
DenseLayer	2048	2048	100.00	100.00	100.00	100.00	261	261
BatchNormLayer	2048	2048	100.00	100.00	100.00	100.00	261	261
FeaturePoolLayer	256	256	100.00	100.00	100.00	100.00	261	261
DropoutLayer	256	256	100.00	100.00	100.00	100.00	261	261
DenseLayer	2048	2048	100.00	100.00	100.00	100.00	261	261
BatchNormLayer	2048	2048	100.00	100.00	100.00	100.00	261	261
FeaturePoolLayer	256	256	100.00	100.00	100.00	100.00	261	261
DenseLayer	10	10	100.00	100.00	100.00	100.00	261	261

Explanation

X, Y: image dimensions
 cap.: learning capacity
 cov.: coverage of image
 magenta: capacity too low (<1/6)
 cyan: image coverage too high (>100%)
 red: capacity too low and coverage too high

epoch	train loss	valid loss	train/val	valid acc	dur
1	35.70711	33.73341	1.05851	0.10000	0.98s
2016-07-31 23:07:59,169	-	root - INFO	- Epoch: 1,	Training error: 35.7071113586	
2	35.62210	33.73750	1.05586	0.10000	0.97s
2016-07-31 23:08:00,149	-	root - INFO	- Epoch: 2,	Training error: 35.6221046448	
3	35.18814	33.73967	1.04293	0.16667	0.98s
2016-07-31 23:08:01,145	-	root - INFO	- Epoch: 3,	Training error: 35.1881446838	
4	34.92326	33.74517	1.03491	0.10000	0.95s
2016-07-31 23:08:02,107	-	root - INFO	- Epoch: 4,	Training error: 34.9232635498	
5	34.53671	33.73351	1.02381	0.06667	0.95s
2016-07-31 23:08:03,070	-	root - INFO	- Epoch: 5,	Training error: 34.5367088318	
6	34.31693	33.72486	1.01756	0.16667	0.95s
2016-07-31 23:08:04,051	-	root - INFO	- Epoch: 6,	Training error: 34.3169288635	
7	34.37460	33.72856	1.01915	0.16667	1.01s
2016-07-31 23:08:05,072	-	root - INFO	- Epoch: 7,	Training error: 34.3745956421	
8	34.40658	33.73640	1.01987	0.13333	1.07s
2016-07-31 23:08:06,168	-	root - INFO	- Epoch: 8,	Training error: 34.406578064	
9	34.21739	33.74545	1.01399	0.06667	1.11s
2016-07-31 23:08:07,293	-	root - INFO	- Epoch: 9,	Training error: 34.2173919678	
10	34.22769	33.73416	1.01463	0.03333	1.00s
2016-07-31 23:08:08,318	-	root - INFO	- Epoch: 10,	Training error: 34.2276878357	
11	34.32386	33.76482	1.01656	0.00000	1.13s
2016-07-31 23:08:09,463	-	root - INFO	- Epoch: 11,	Training error: 34.3238601685	
12	34.36395	33.89575	1.01381	0.10000	1.11s
2016-07-31 23:08:10,584	-	root - INFO	- Epoch: 12,	Training error: 34.3639450073	
13	34.49736	34.00295	1.01454	0.10000	1.05s
2016-07-31 23:08:11,643	-	root - INFO	- Epoch: 13,	Training error: 34.4973640442	
14	34.40776	34.89201	0.98612	0.06667	1.01s
2016-07-31 23:08:12,658	-	root - INFO	- Epoch: 14,	Training error: 34.4077644348	
15	34.62708	36.87984	0.93892	0.06667	0.94s
2016-07-31 23:08:13,605	-	root - INFO	- Epoch: 15,	Training error: 34.62707901	
16	34.87041	38.94115	0.89546	0.13333	0.98s
2016-07-31 23:08:14,589	-	root - INFO	- Epoch: 16,	Training error: 34.8704109192	
17	34.80553	42.91374	0.81106	0.13333	0.97s
2016-07-31 23:08:15,575	-	root - INFO	- Epoch: 17,	Training error: 34.8055305481	
18	34.65967	45.42543	0.76300	0.10000	0.99s
2016-07-31 23:08:16,578	-	root - INFO	- Epoch: 18,	Training error: 34.6596679688	
19	34.74032	47.51985	0.73107	0.10000	0.93s
2016-07-31 23:08:17,522	-	root - INFO	- Epoch: 19,	Training error: 34.740322113	
20	34.52382	48.51399	0.71163	0.10000	0.96s
2016-07-31 23:08:18,492	-	root - INFO	- Epoch: 20,	Training error: 34.5238227844	
21	34.68174	51.85040	0.66888	0.06667	0.98s
2016-07-31 23:08:19,476	-	root - INFO	- Epoch: 21,	Training error: 34.6817359924	
22	34.51344	52.82178	0.65339	0.10000	1.10s
2016-07-31 23:08:20,601	-	root - INFO	- Epoch: 22,	Training error: 34.5134353638	
23	34.35563	51.70115	0.66450	0.13333	1.00s
2016-07-31 23:08:21,620	-	root - INFO	- Epoch: 23,	Training error: 34.355632782	
24	34.32787	51.50606	0.66648	0.06667	0.96s
2016-07-31 23:08:22,594	-	root - INFO	- Epoch: 24,	Training error: 34.3278694153	
25	34.24430	48.83182	0.70127	0.10000	0.97s
2016-07-31 23:08:23,577	-	root - INFO	- Epoch: 25,	Training error: 34.2443008423	

26	34.23243	45.91653	0.74554	0.10000	1.01s
2016-07-31 23:08:24,595	- root - INFO - Epoch: 26, Training error: 34.2324295044				
27	34.08209	48.21726	0.70684	0.13333	0.98s
2016-07-31 23:08:25,587	- root - INFO - Epoch: 27, Training error: 34.0820922852				
28	34.28476	54.34813	0.63084	0.10000	0.95s
2016-07-31 23:08:26,551	- root - INFO - Epoch: 28, Training error: 34.2847557068				
29	33.99178	66.22639	0.51327	0.10000	1.00s
2016-07-31 23:08:27,561	- root - INFO - Epoch: 29, Training error: 33.9917793274				
30	33.98634	71.18796	0.47742	0.10000	0.98s
2016-07-31 23:08:28,548	- root - INFO - Epoch: 30, Training error: 33.9863357544				
31	34.20824	94.85815	0.36063	0.10000	0.97s
2016-07-31 23:08:29,530	- root - INFO - Epoch: 31, Training error: 34.208240509				
32	34.08863	97.63881	0.34913	0.10000	0.92s
2016-07-31 23:08:30,460	- root - INFO - Epoch: 32, Training error: 34.088634491				
33	34.13584	100.28867	0.34038	0.10000	0.95s
2016-07-31 23:08:31,420	- root - INFO - Epoch: 33, Training error: 34.1358413696				
34	34.07228	99.97665	0.34080	0.10000	0.99s
2016-07-31 23:08:32,437	- root - INFO - Epoch: 34, Training error: 34.0722808838				
35	34.13663	106.77155	0.31972	0.10000	0.99s
2016-07-31 23:08:33,448	- root - INFO - Epoch: 35, Training error: 34.1366348267				
36	34.22134	135.58780	0.25239	0.10000	0.97s
2016-07-31 23:08:34,433	- root - INFO - Epoch: 36, Training error: 34.2213439941				
37	34.27357	144.03143	0.23796	0.10000	1.03s
2016-07-31 23:08:35,475	- root - INFO - Epoch: 37, Training error: 34.2735710144				
38	34.06572	135.94006	0.25059	0.10000	0.97s
2016-07-31 23:08:36,464	- root - INFO - Epoch: 38, Training error: 34.0657196045				
39	34.10125	145.17856	0.23489	0.10000	1.03s
2016-07-31 23:08:37,508	- root - INFO - Epoch: 39, Training error: 34.1012458801				
40	34.01466	124.50072	0.27321	0.10000	0.98s
2016-07-31 23:08:38,503	- root - INFO - Epoch: 40, Training error: 34.0146598816				

2.2 Visualizations

```
In [ ]: from notebook_functions import plot_validation_loss
```

```
In [ ]: plot_validation_loss(net, validation_file_name, ylim=[0, 0.5])
```

```
In [ ]:
```