

run_train_distracted_drivers-maxout8

August 4, 2016

The changes here are: - Orthogonal weight initialization - Decreasing learning rate by step size

```
In [1]: from skimage import io, transform, exposure, color, util
import os, itertools, sys
from PIL import Image
%pylab inline
sys.setrecursionlimit(1000000)
```

Populating the interactive namespace from numpy and matplotlib

```
In [2]: # data_dir = "/home/dylan/IdeaProjects/distracted_drivers/train/"
data_dir = "/media/dylan/Science/Kaggle-Data/distracted_drivers/train/"
```

```
In [3]: input_volume_shape = (128, 128)
```

```
In [4]: def read_img_file_PIL(file_path, size=(32,32)):
    img = Image.open(file_path).convert('L')
    img.thumbnail(size, Image.NEAREST)
    data = np.array(img)
    shape = data.shape
    append_top = int(ceil(max(0, size[0] - shape[0])/2.0))
    append_bot = int(floor(max(0, size[0] - shape[0])/2.0))
    data = util.pad(data, ((append_top, append_bot),
                           (0,0)), mode='constant', constant_values=0)

    return data
```

```
In [5]: def read_img_file(file_path, rescale=0.01):
    img = io.imread(file_path)
    img = color.rgb2gray(img)
    return transform.rescale(img, rescale)
```

```
In [6]: def image_gen_from_dir(directory, batch_size, num_categories, size=input_volume_shape):
    result = {os.path.join(dp, f) : int(os.path.split(dp)[1]) for dp, dn, filenames in os.walk(
        directory)
        for f in filenames if os.path.splitext(f)[1] == '.jpg'}

    # infinite loop
    while True:
        image_files = []
        labels = []
        # randomly choose batch size samples in result
        for category in range(num_categories):
            file_samples = np.random.choice([k for k, v in result.iteritems() if v == category],
                                             size=batch_size, replace=False)
            for file_sample in file_samples:
                image_files.append(read_img_file_PIL(file_sample, size=size))
            labels.extend([category for v in itertools.repeat(category, batch_size)])
```

```

        # end category loop
    X = np.asarray(image_files, dtype=np.float32)
    # -1 to 1 range
    X = exposure.rescale_intensity(X, out_range=(-1,1))
    y = np.asarray(labels, dtype=np.int32)
    yield X, y

```

0.1 Another loader, augmentation time

We'll do 6 augmentations:

- 1.) Translation up to 10 pixels
- 2.) Rotation up to 15 degrees
- 3.) Zooming
- 4.) JPEG compression
- 5.) Sharpening
- 6.) Gamma correction

We won't do flips since the dataset only contains images from the passenger seat. Perhaps we can revisit this later.

```

In [7]: from skimage.transform import rotate, warp, AffineTransform
        from skimage import filters
        from scipy import ndimage, misc
        import StringIO

```

```

In [8]: def random_translate(img):
        shift_random = AffineTransform(translation=(randint(-10, 10), randint(-10, 10)))
        min_value = 0 if min(img.ravel()) > 0 else min(img.ravel())
        return np.float32(warp(img, shift_random, mode='constant', cval=min_value))

```

```

def random_rotate(img):
    min_value = 0 if min(img.ravel()) > 0 else min(img.ravel())
    return np.float32(rotate(img, randint(-15, 15), mode='constant', cval=min_value))

```

```

def random_zoom(img):
    min_value = 0 if min(img.ravel()) > 0 else min(img.ravel())
    scale_random = AffineTransform(scale=(uniform(0.9, 1.1), uniform(0.9, 1.1)))
    return np.float32(warp(img, scale_random, mode='constant', cval=min_value))

```

```

def random_compress(img):
    max_v = np.ceil(img.max())
    min_v = np.floor(img.min())
    nd_im = exposure.rescale_intensity(img, out_range=(0, 1)).squeeze()
    nd_im = np.ndarray.astype(nd_im * 255, np.uint8)
    # nd_im = np.ndarray.astype(img * 255, np.uint8)
    im = Image.fromarray(nd_im)
    buf = StringIO.StringIO()
    im.save(buf, "JPEG", quality=np.random.randint(95, 99))
    buf.seek(0)
    im2 = Image.open(buf)
    x1 = exposure.rescale_intensity(np.ndarray.astype(np.array(im2), np.float32), out_range=(min
    return x1

```

```

def random_sharpening(img):
    blurred_f = ndimage.gaussian_filter(img, 0.5)
    filter_blurred_f = ndimage.gaussian_filter(blurred_f, 1)
    alpha = uniform(0.9, 1.2)
    img = blurred_f + alpha * (blurred_f - filter_blurred_f)
    return exposure.rescale_intensity(img, out_range=(-1, 1))

def random_gamma_correction(img):
    max_v = np.ceil(img.max())
    min_v = np.floor(img.min())
    img = exposure.rescale_intensity(img, out_range=(0,1))
    img = exposure.adjust_gamma(img, uniform(0.2, 0.8))
    return exposure.rescale_intensity(img, out_range=(-1, 1))

In [9]: def random_aug(img):
    choice = np.random.randint(0,6)
    # choose from 4 different augmentations!
    if choice == 0:
        return random_translate(img)
    elif choice == 1:
        return random_rotate(img)
    elif choice == 2:
        return random_zoom(img)
    elif choice == 3:
        return random_compress(img)
    elif choice == 4:
        return random_sharpening(img)
    else:
        return random_gamma_correction(img)

In [10]: def random_aug_batch(X, aug_algorithm):
    for i in range(X.shape[0]):
        X[i] = aug_algorithm(X[i])
    return X

In [11]: def random_aug_gen(gen, aug_algorithm):
    for batchX, batchY in gen:
        yield random_aug_batch(batchX, aug_algorithm), batchY

```

1 Process Generator with cached elements

```

In [12]: def threaded_generator(generator, num_cached=50):
    import Queue
    queue = Queue.Queue(maxsize=num_cached)
    sentinel = object() # guaranteed unique reference

    # define producer (putting items into queue)
    def producer():
        for item in generator:
            queue.put(item)
        queue.put(sentinel)

    # start producer (in a background thread)
    import threading

```

```

thread = threading.Thread(target=producer)
thread.daemon = True
thread.start()

# run as consumer (read items from queue, in current thread)
item = queue.get()
while item is not sentinel:
    yield item
    queue.task_done()
    item = queue.get()

```

```

In [13]: from nolearn.lasagne import NeuralNet
        from lasagne.layers import DenseLayer, ReshapeLayer, Upscale2DLayer, Conv2DLayer, InputLayer,
        MaxPool2DLayer, get_all_params, batch_norm, BatchNormLayer, FeaturePoolLayer
        import numpy as np
        from lasagne.nonlinearities import softmax, leaky_rectify, theano
        from lasagne.updates import nesterov_momentum
        from nolearn.lasagne import NeuralNet, BatchIterator, PrintLayerInfo, objective
        from nolearn.lasagne import TrainSplit
        from common import EarlyStopping, EndTrainingFromEarlyStopping
        from lasagne.objectives import categorical_crossentropy, aggregate
        import cPickle as pickle
        from sklearn import metrics
        import time, logging, logging.config, logging.handlers
        from lasagne.init import Orthogonal
        from notebook_functions import load_best_weights

```

Couldn't import dot_parser, loading of dot files will not be possible.

Using gpu device 0: GeForce GTX 960 (CNMeM is disabled, CuDNN 4004)

```
def batch_norm(s): return s
```

In [14]: try:

```

from lasagne.layers.dnn import Conv2DDNNLayer, MaxPool2DDNNLayer
def conv_2_layer_stack(top, num_filters):
    conv1 = batch_norm(Conv2DDNNLayer(top, num_filters, (3, 3),
        stride=1, pad=1, nonlinearity=leaky_rectify, W=Orthogonal()))
    conv2 = batch_norm(Conv2DDNNLayer(conv1, num_filters, (3, 3),
        stride=1, pad=1, nonlinearity=leaky_rectify, W=Orthogonal()))
    return MaxPool2DDNNLayer(conv2, (2, 2), 2)

def conv_3_layer_stack(top, num_filters):
    conv1 = batch_norm(Conv2DDNNLayer(top, num_filters, (3, 3),
        stride=1, pad=1, nonlinearity=leaky_rectify, W=Orthogonal()))
    conv2 = batch_norm(Conv2DDNNLayer(conv1, num_filters, (3, 3),
        stride=1, pad=1, nonlinearity=leaky_rectify, W=Orthogonal()))
    conv3 = batch_norm(Conv2DDNNLayer(conv2, num_filters, (3, 3),
        stride=1, pad=1, nonlinearity=leaky_rectify, W=Orthogonal()))
    return MaxPool2DDNNLayer(conv3, (2, 2), 2)

def conv_4_layer_stack(top, num_filters):
    conv1 = batch_norm(Conv2DDNNLayer(top, num_filters, (3, 3),
        stride=1, pad=1, nonlinearity=leaky_rectify, W=Orthogonal()))
    conv2 = batch_norm(Conv2DDNNLayer(conv1, num_filters, (3, 3),

```

```

        stride=1, pad=1, nonlinearity=leaky_rectify, W=Orthogonal()))
conv3 = batch_norm(Conv2DDNNLayer(conv2, num_filters, (3, 3),
        stride=1, pad=1, nonlinearity=leaky_rectify, W=Orthogonal()))
conv4 = batch_norm(Conv2DDNNLayer(conv3, num_filters, (3, 3),
        stride=1, pad=1, nonlinearity=leaky_rectify, W=Orthogonal()))
return MaxPool2DDNNLayer(conv4, (2, 2), 2)

def conv_6_layer_stack(top, num_filters):
conv1 = batch_norm(Conv2DDNNLayer(top, num_filters, (3, 3),
        stride=1, pad=1, nonlinearity=leaky_rectify, W=Orthogonal()))
conv2 = batch_norm(Conv2DDNNLayer(conv1, num_filters, (3, 3),
        stride=1, pad=1, nonlinearity=leaky_rectify, W=Orthogonal()))
conv3 = batch_norm(Conv2DDNNLayer(conv2, num_filters, (3, 3),
        stride=1, pad=1, nonlinearity=leaky_rectify, W=Orthogonal()))
conv4 = batch_norm(Conv2DDNNLayer(conv3, num_filters, (3, 3),
        stride=1, pad=1, nonlinearity=leaky_rectify, W=Orthogonal()))
conv5 = batch_norm(Conv2DDNNLayer(conv4, num_filters, (3, 3),
        stride=1, pad=1, nonlinearity=leaky_rectify, W=Orthogonal()))
conv6 = batch_norm(Conv2DDNNLayer(conv5, num_filters, (3, 3),
        stride=1, pad=1, nonlinearity=leaky_rectify, W=Orthogonal()))
return MaxPool2DLayer(conv6, (2, 2), 2)

except ImportError:
def conv_2_layer_stack(top, num_filters):
conv1 = batch_norm(Conv2DLayer(top, num_filters, (3, 3), stride=1, pad=1, nonlinearity=
conv2 = batch_norm(Conv2DLayer(conv1, num_filters, (3, 3), stride=1, pad=1, nonlineari
return MaxPool2DLayer(conv2, (2, 2), 2)

def conv_3_layer_stack(top, num_filters):
conv1 = batch_norm(Conv2DLayer(top, num_filters, (3, 3), stride=1, pad=1, nonlinearity=
conv2 = batch_norm(Conv2DLayer(conv1, num_filters, (3, 3), stride=1, pad=1, nonlineari
conv3 = batch_norm(Conv2DLayer(conv2, num_filters, (3, 3), stride=1, pad=1, nonlineari
return MaxPool2DLayer(conv3, (2, 2), 2)

def conv_4_layer_stack(top, num_filters):
conv1 = batch_norm(Conv2DLayer(top, num_filters, (3, 3), stride=1, pad=1, nonlinearity=
conv2 = batch_norm(Conv2DLayer(conv1, num_filters, (3, 3), stride=1, pad=1, nonlineari
conv3 = batch_norm(Conv2DLayer(conv2, num_filters, (3, 3), stride=1, pad=1, nonlineari
conv4 = batch_norm(Conv2DLayer(conv3, num_filters, (3, 3), stride=1, pad=1, nonlineari
return MaxPool2DLayer(conv4, (2, 2), 2)

def conv_6_layer_stack(top, num_filters):
conv1 = batch_norm(Conv2DLayer(top, num_filters, (3, 3), stride=1, pad=1, nonlinearity=
conv2 = batch_norm(Conv2DLayer(conv1, num_filters, (3, 3), stride=1, pad=1, nonlineari
conv3 = batch_norm(Conv2DLayer(conv2, num_filters, (3, 3), stride=1, pad=1, nonlineari
conv4 = batch_norm(Conv2DLayer(conv3, num_filters, (3, 3), stride=1, pad=1, nonlineari
conv5 = batch_norm(Conv2DLayer(conv4, num_filters, (3, 3), stride=1, pad=1, nonlineari
conv6 = batch_norm(Conv2DLayer(conv5, num_filters, (3, 3), stride=1, pad=1, nonlineari
return MaxPool2DLayer(conv6, (2, 2), 2)

In [15]: k = 8
input_layer = InputLayer((None, 1, input_volume_shape[0], input_volume_shape[1]))
conv_stack_1 = conv_2_layer_stack(input_layer, 32)
dropout1 = DropoutLayer(conv_stack_1, p=0.1)

```

```

conv_stack_2 = conv_2_layer_stack(dropout1, 64)
dropout2 = DropoutLayer(conv_stack_2, p=0.2)

conv_stack_3 = conv_2_layer_stack(dropout2, 128)
dropout3 = DropoutLayer(conv_stack_3, p=0.3)

conv_stack_4 = conv_2_layer_stack(dropout3, 256)
dropout4 = DropoutLayer(conv_stack_4, p=0.4)

conv_stack_5 = conv_2_layer_stack(dropout4, 512)
dropout17 = DropoutLayer(conv_stack_5, p=0.5)

dense18 = DenseLayer(dropout17, 2048, nonlinearity=None)
norm1 = BatchNormLayer(dense18)
maxout1 = FeaturePoolLayer(norm1, k)
dropout19 = DropoutLayer(maxout1, p=0.5)

dense20 = DenseLayer(dropout19, 2048, nonlinearity=None)
norm2 = BatchNormLayer(dense20)
maxout2 = FeaturePoolLayer(norm2, k)

softmax21 = DenseLayer(maxout2, 10, nonlinearity=softmax)

```

1.1 Quality of Life Functions

```

In [16]: if not os.path.exists("logs"):
        os.mkdir("logs")
        logging.config.fileConfig("logging-training.conf")

def regularization_objective(layers, lambda1=0., lambda2=0., *args, **kwargs):
    # default loss
    losses = objective(layers, *args, **kwargs)
    # get layer weights except for the biases
    weights = get_all_params(layers[-1], regularizable=True)
    regularization_term = 0.0
    # sum of abs weights for L1 regularization
    if lambda1 != 0.0:
        sum_abs_weights = sum([abs(w).sum() for w in weights])
        regularization_term += (lambda1 * sum_abs_weights)
    # sum of squares (sum(theta^2))
    if lambda2 != 0.0:
        sum_squared_weights = (1 / 2.0) * sum([(w ** 2).sum() for w in weights])
        regularization_term += (lambda2 * sum_squared_weights)
    # add weights to regular loss
    losses += regularization_term
    return losses

def eval_regularization(net):
    if net.objective_lambda1 == 0 and net.objective_lambda2 == 0:
        return 0
    # check the loss if the regularization term is not overpowering the loss
    weights = get_all_params(net.layers[-1], regularizable=True)
    # sum of abs weights for L1 regularization

```

```

sum_abs_weights = sum([abs(w).sum() for w in weights])
# sum of squares (sum(theta^2))
sum_squared_weights = (1 / 2.0) * sum([(w ** 2).sum() for w in weights])
# add weights to regular loss
regularization_term = (net.objective_lambda1 * sum_abs_weights) \
    + (net.objective_lambda2 * sum_squared_weights)
return regularization_term

def print_regularization_term(net):
    if net.objective_lambda1 > 0.0 or net.objective_lambda2 > 0.0:
        regularization_term = eval_regularization(net)
        print "Regularization term: {}".format(regularization_term.eval())

def validation_set_loss(_net, _X, _y):
    """We need this to track the validation loss"""
    _yb = _net.predict_proba(_X)
    _y_pred = np.argmax(_yb, axis=1)
    _acc = metrics.accuracy_score(_y, _y_pred)
    loss = aggregate(categorical_crossentropy(_yb, _y))
    loss += eval_regularization(_net)
    return loss, _acc

def store_model(model_file_name, net):
    directory_name = os.path.dirname(model_file_name)
    model_file_name = os.path.basename(model_file_name)
    if not os.path.exists(directory_name):
        os.makedirs(directory_name)
    # write model
    output_model_file_name = os.path.join(directory_name, model_file_name)
    start_write_time = time.time()
    if os.path.isfile(output_model_file_name):
        os.remove(output_model_file_name)
    with open(output_model_file_name, 'wb') as experiment_model:
        pickle.dump(net, experiment_model)
    total_write_time = time.time() - start_write_time
    m, s = divmod(total_write_time, 60)
    h, m = divmod(m, 60)
    logging.log(logging.INFO, "Duration of saving to disk: %0d:%02d:%02d", h, m, s)

def write_validation_loss_and_store_best(validation_file_name, best_weights_file_name,
                                         net, X_val, y_val, best_vloss, best_acc):
    # write validation loss
    start_validate_time = time.time()
    vLoss, vAcc = validation_set_loss(net, X_val, y_val)
    loss = vLoss.eval()
    current_epoch = net.train_history_[-1]['epoch']
    with open(validation_file_name, 'a') as validation_file:
        validation_file.write("{} {}, {}{}\n".format(current_epoch, loss, vAcc))

    total_validate_time = time.time() - start_validate_time
    m, s = divmod(total_validate_time, 60)
    h, m = divmod(m, 60)

```

```

logging.log(logging.INFO, "Duration of validation: %0d:%02d:%02d", h, m, s)

# store best weights here
if loss < best_vloss:
    start_bw_time = time.time()
    best_vloss = loss
    best_acc = vAcc
    with open(best_weights_file_name, 'wb') as best_model_file:
        pickle.dump(net.get_all_params_values(), best_model_file, -1)

return best_vloss, best_acc

class AdjustVariableWithStepSize(object):
    """This class adjusts any variable during training"""

    def __init__(self, name, start=0.03, steps=3, after_epochs=2000):
        self.name = name
        self.start = start
        self.steps=steps
        self.after_epochs=after_epochs
        self.ls = []

    def __call__(self, nn, train_history):
        if not self.ls:
            for i in range(self.steps):
                self.ls.extend(np.repeat(self.start/(np.power(10,i)), self.after_epochs))

        try:
            epoch = train_history[-1]['epoch']
            new_value = np.float32(self.ls[epoch - 1])
            getattr(nn, self.name).set_value(new_value)
        except IndexError:
            pass

```

1.2 CNN

```

In [17]: lambda1 = 0.0
         lambda2 = 5e-3

```

```

net = NeuralNet(
    layers=softmax21,
    max_epochs=1,
    update=nesterov_momentum,
    update_learning_rate=theano.shared(np.float32(0.001)),
    update_momentum = 0.99,
    # update=adam,
    on_epoch_finished=[
        EarlyStopping(patience=1000),
        AdjustVariableWithStepSize('update_learning_rate', start=0.001, steps=2, after_epochs=
    ],
    on_training_finished=[
        EndTrainingFromEarlyStopping()

```



```

],
objective=regularization_objective,
objective_lambda2=lambda2,
objective_lambda1=lambda1,
batch_iterator_train=BatchIterator(batch_size=100),
train_split=TrainSplit(
    eval_size=0.25),
# train_split=TrainSplit(eval_size=0.0),
verbose=3,
)

```

```

In [18]: p = PrintLayerInfo()
net.initialize()
# p(net)

```

1.2.1 load cnn instead

```

In [19]: dir_name = 'net.vgg.large.12.5e3.orthog-norm-maxout8'
validation_file_name = "{}/vloss-{}.txt".format(dir_name, dir_name)
model_file_name = "{}/{}.pickle".format(dir_name, dir_name)
best_weights_file_name = "{}/bw-{}.weights".format(dir_name, dir_name)
if os.path.exists(dir_name):
    print "Model exists. Loading {}".format(dir_name)
    with open(model_file_name, 'rb') as reader:
        net = pickle.load(reader)
else:
    print "Training model from the beginning {}".format(dir_name)

```

Model exists. Loading net.vgg.large.12.5e3.orthog-norm-maxout8.

```

In [20]: load_best_weights(best_weights_file_name, net)

```

```

Loaded parameters to layer 'conv2ddnn1' (shape 32x1x3x3).
Loaded parameters to layer 'batchnorm2' (shape 32).
Loaded parameters to layer 'batchnorm2' (shape 32).
Loaded parameters to layer 'batchnorm2' (shape 32).
Loaded parameters to layer 'batchnorm2' (shape 32).
Loaded parameters to layer 'conv2ddnn4' (shape 32x32x3x3).
Loaded parameters to layer 'batchnorm5' (shape 32).
Loaded parameters to layer 'batchnorm5' (shape 32).
Loaded parameters to layer 'batchnorm5' (shape 32).
Loaded parameters to layer 'batchnorm5' (shape 32).
Loaded parameters to layer 'conv2ddnn9' (shape 64x32x3x3).
Loaded parameters to layer 'batchnorm10' (shape 64).
Loaded parameters to layer 'batchnorm10' (shape 64).
Loaded parameters to layer 'batchnorm10' (shape 64).
Loaded parameters to layer 'batchnorm10' (shape 64).
Loaded parameters to layer 'conv2ddnn12' (shape 64x64x3x3).
Loaded parameters to layer 'batchnorm13' (shape 64).
Loaded parameters to layer 'batchnorm13' (shape 64).
Loaded parameters to layer 'batchnorm13' (shape 64).
Loaded parameters to layer 'batchnorm13' (shape 64).
Loaded parameters to layer 'conv2ddnn17' (shape 128x64x3x3).
Loaded parameters to layer 'batchnorm18' (shape 128).
Loaded parameters to layer 'batchnorm18' (shape 128).

```

```

Loaded parameters to layer 'batchnorm18' (shape 128).
Loaded parameters to layer 'batchnorm18' (shape 128).
Loaded parameters to layer 'conv2ddnn20' (shape 128x128x3x3).
Loaded parameters to layer 'batchnorm21' (shape 128).
Loaded parameters to layer 'batchnorm21' (shape 128).
Loaded parameters to layer 'batchnorm21' (shape 128).
Loaded parameters to layer 'batchnorm21' (shape 128).
Loaded parameters to layer 'conv2ddnn25' (shape 256x128x3x3).
Loaded parameters to layer 'batchnorm26' (shape 256).
Loaded parameters to layer 'batchnorm26' (shape 256).
Loaded parameters to layer 'batchnorm26' (shape 256).
Loaded parameters to layer 'batchnorm26' (shape 256).
Loaded parameters to layer 'conv2ddnn28' (shape 256x256x3x3).
Loaded parameters to layer 'batchnorm29' (shape 256).
Loaded parameters to layer 'batchnorm29' (shape 256).
Loaded parameters to layer 'batchnorm29' (shape 256).
Loaded parameters to layer 'conv2ddnn33' (shape 512x256x3x3).
Loaded parameters to layer 'batchnorm34' (shape 512).
Loaded parameters to layer 'batchnorm34' (shape 512).
Loaded parameters to layer 'batchnorm34' (shape 512).
Loaded parameters to layer 'batchnorm34' (shape 512).
Loaded parameters to layer 'conv2ddnn36' (shape 512x512x3x3).
Loaded parameters to layer 'batchnorm37' (shape 512).
Loaded parameters to layer 'batchnorm37' (shape 512).
Loaded parameters to layer 'batchnorm37' (shape 512).
Loaded parameters to layer 'batchnorm37' (shape 512).
Loaded parameters to layer 'dense41' (shape 8192x2048).
Loaded parameters to layer 'dense41' (shape 2048).
Loaded parameters to layer 'batchnorm42' (shape 2048).
Loaded parameters to layer 'batchnorm42' (shape 2048).
Loaded parameters to layer 'batchnorm42' (shape 2048).
Loaded parameters to layer 'batchnorm42' (shape 2048).
Loaded parameters to layer 'dense45' (shape 256x2048).
Loaded parameters to layer 'dense45' (shape 2048).
Loaded parameters to layer 'batchnorm46' (shape 2048).
Loaded parameters to layer 'batchnorm46' (shape 2048).
Loaded parameters to layer 'batchnorm46' (shape 2048).
Loaded parameters to layer 'batchnorm46' (shape 2048).
Loaded parameters to layer 'dense48' (shape 256x10).
Loaded parameters to layer 'dense48' (shape 10).

```

```


In [21]: from nolearn.lasagne.visualize import plot_loss
         plt.figure( figsize=(15,9))
         plt.ylim([0.1,0.5])
         plt.plot([v['valid_loss'] for v in net.train_history_])

```

```

Out[21]: [<matplotlib.lines.Line2D at 0x7f8fcd494b10>]

```



```
In [22]: net.on_epoch_finished
```

```
Out[22]: [<nolearn.lasagne.handlers.PrintLog instance at 0x7f8fcabb9248>]
```

```
In [23]: net.update_learning_rate = 0.00001
         # net.on_epoch_finished.pop(0)
```

2 just this time.

```
net.on_epoch_finished.pop(1) print net.on_epoch_finished net.update_learning_rate=0.001
```

2.1 Define validation set

```
In [ ]: val_dir = "/media/dylan/Science/Kaggle-Data/distracted_drivers/val/"
        X_val, y_val = image_gen_from_dir(val_dir, 40, 10, size=input_volume_shape).next()
        X_val = X_val.reshape(-1, 1, input_volume_shape[0], input_volume_shape[1])

In [ ]: image_gen = image_gen_from_dir(data_dir, 10, 10, size=input_volume_shape)
        gen = random_aug_gen(image_gen, random_aug)
        threaded_gen = threaded_generator(gen, num_cached=100)

ops_every = 500
best_acc = 0.0
best_vloss = np.inf

start_time = time.time()
try:
    for step, (inputs, targets) in enumerate(threaded_gen):
        shape = inputs.shape
        net.fit(inputs.reshape(shape[0],1, shape[1], shape[2]), targets)
```

```

        if (step + 1) % ops_every == 0:
            print_regularization_term(net)
            store_model(model_file_name, net)
            # center validation
            best_vloss, best_acc = write_validation_loss_and_store_best(
                validation_file_name, best_weights_file_name, net, X_val, y_val, best_vloss, be

except StopIteration:
    # terminate if already early stopping
    with open(model_file_name, 'wb') as writer:
        pickle.dump(net, writer)
    total_time = time.time() - start_time
    print("Training successful by early stopping. Elapsed: {}".format(total_time))

```

53903	0.54187	0.56071	0.96640	0.96667	1.04s
53904	0.54011	0.67385	0.80153	0.96667	1.02s
53905	0.57318	0.49750	1.15213	1.00000	1.01s
53906	0.55604	0.49808	1.11635	1.00000	1.02s
53907	0.53049	0.54025	0.98195	1.00000	1.03s
53908	0.53793	0.49684	1.08270	1.00000	1.04s
53909	0.58136	0.53941	1.07777	0.96667	1.03s
53910	0.52917	0.50147	1.05524	1.00000	1.03s
53911	0.54986	0.49550	1.10971	1.00000	1.03s
53912	0.60367	0.52409	1.15183	0.96667	1.04s
53913	0.56180	0.51511	1.09063	1.00000	1.02s
53914	0.55625	0.49974	1.11308	1.00000	1.05s
53915	0.54149	0.56290	0.96196	0.96667	1.03s
53916	0.53287	0.49828	1.06942	1.00000	1.03s
53917	0.51200	0.50430	1.01527	1.00000	1.03s
53918	0.51252	0.77564	0.66077	0.93333	1.02s
53919	0.53798	0.60664	0.88681	0.93333	1.03s

2.2 Visualizations

```

In [ ]: from notebook_functions import plot_validation_loss

In [ ]: plot_validation_loss(net, validation_file_name, ylim=[0, 0.5])

In [ ]:

```