

# run\_train\_distracted\_drivers.py

August 4, 2016

```
In [1]: from skimage import io, transform, exposure, color, util
import os, itertools, sys
from PIL import Image
%pylab inline
sys.setrecursionlimit(1000000)
```

Populating the interactive namespace from numpy and matplotlib

```
In [2]: # data_dir = "/home/dylan/IdeaProjects/distracted_drivers/train/"
data_dir = "/media/dylan/Science/Kaggle-Data/distracted_drivers/train/"
```

```
In [3]: input_volume_shape = (128, 128)
```

```
In [4]: def read_img_file_PIL(file_path, size=(32,32)):
    img = Image.open(file_path).convert('L')
    img.thumbnail(size, Image.NEAREST)
    data = np.array(img)
    shape = data.shape
    append_top = int(ceil(max(0, size[0] - shape[0])/2.0))
    append_bot = int(floor(max(0, size[0] - shape[0])/2.0))
    data = util.pad(data, ((append_top, append_bot),
                           (0,0)), mode='constant', constant_values=0)

    return data
```

```
In [5]: def read_img_file(file_path, rescale=0.01):
    img = io.imread(file_path)
    img = color.rgb2gray(img)
    return transform.rescale(img, rescale)
```

```
In [6]: def image_gen_from_dir(directory, batch_size, num_categories, size=input_volume_shape):
    result = {os.path.join(dp, f) : int(os.path.split(dp)[1]) for dp, dn, filenames in os.walk(directory)
              for f in filenames if os.path.splitext(f)[1] == '.jpg'}

    # infinite loop
    while True:
        image_files = []
        labels = []
        # randomly choose batch size samples in result
        for category in range(num_categories):
            file_samples = np.random.choice([k for k, v in result.iteritems() if v == category],
                                             size=batch_size, replace=False)
            for file_sample in file_samples:
                image_files.append(read_img_file_PIL(file_sample, size=size))
                labels.extend([v for v in itertools.repeat(category, batch_size)])

    # end category loop
```

```

X = np.asarray(image_files, dtype=np.float32)
# -1 to 1 range
X = exposure.rescale_intensity(X, out_range=(-1,1))
y = np.asarray(labels, dtype=np.int32)
yield X, y

```

## 0.1 Another loader, augmentation time

We'll do 6 augmentations:

- 1.) Translation up to 10 pixels
- 2.) Rotation up to 15 degrees
- 3.) Zooming
- 4.) JPEG compression
- 5.) Sharpening
- 6.) Gamma correction

We won't do flips since the dataset only contains images from the passenger seat. Perhaps we can revisit this later.

```

In [7]: from skimage.transform import rotate, warp, AffineTransform
        from skimage import filters
        from scipy import ndimage, misc
        import StringIO

```

```

In [8]: def random_translate(img):
        shift_random = AffineTransform(translation=(randint(-10, 10), randint(-10, 10)))
        min_value = 0 if min(img.ravel()) > 0 else min(img.ravel())
        return np.float32(warp(img, shift_random, mode='constant', cval=min_value))

```

```

def random_rotate(img):
    min_value = 0 if min(img.ravel()) > 0 else min(img.ravel())
    return np.float32(rotate(img, randint(-15, 15), mode='constant', cval=min_value))

```

```

def random_zoom(img):
    min_value = 0 if min(img.ravel()) > 0 else min(img.ravel())
    scale_random = AffineTransform(scale=(uniform(0.9, 1.1), uniform(0.9, 1.1)))
    return np.float32(warp(img, scale_random, mode='constant', cval=min_value))

```

```

def random_compress(img):
    max_v = np.ceil(img.max())
    min_v = np.floor(img.min())
    nd_im = exposure.rescale_intensity(img, out_range=(0, 1)).squeeze()
    nd_im = np.ndarray.astype(nd_im * 255, np.uint8)
    # nd_im = np.ndarray.astype(img * 255, np.uint8)
    im = Image.fromarray(nd_im)
    buf = StringIO.StringIO()
    im.save(buf, "JPEG", quality=np.random.randint(95, 99))
    buf.seek(0)
    im2 = Image.open(buf)
    x1 = exposure.rescale_intensity(np.ndarray.astype(np.array(im2), np.float32), out_range=(min_v, max_v))
    return x1

```

```

def random_sharpening(img):
    blurred_f = ndimage.gaussian_filter(img, 0.5)

```

```

filter_blurred_f = ndimage.gaussian_filter(blurred_f, 1)
alpha = uniform(0.9, 1.2)
img = blurred_f + alpha * (blurred_f - filter_blurred_f)
return exposure.rescale_intensity(img, out_range=(-1, 1))

def random_gamma_correction(img):
    max_v = np.ceil(img.max())
    min_v = np.floor(img.min())
    img = exposure.rescale_intensity(img, out_range=(0,1))
    img = exposure.adjust_gamma(img, uniform(0.2, 0.8))
    return exposure.rescale_intensity(img, out_range=(-1, 1))

In [9]: def random_aug(img):
    choice = np.random.randint(0,6)
    # choose from 4 different augmentations!
    if choice == 0:
        return random_translate(img)
    elif choice == 1:
        return random_rotate(img)
    elif choice == 2:
        return random_zoom(img)
    elif choice == 3:
        return random_compress(img)
    elif choice == 4:
        return random_sharpening(img)
    else:
        return random_gamma_correction(img)

In [10]: def random_aug_batch(X, aug_algorithm):
    for i in range(X.shape[0]):
        X[i] = aug_algorithm(X[i])
    return X

In [11]: def random_aug_gen(gen, aug_algorithm):
    for batchX, batchY in gen:
        yield random_aug_batch(batchX, aug_algorithm), batchY

```

## 1 Process Generator with cached elements

```

In [12]: def threaded_generator(generator, num_cached=50):
    import Queue
    queue = Queue.Queue(maxsize=num_cached)
    sentinel = object() # guaranteed unique reference

    # define producer (putting items into queue)
    def producer():
        for item in generator:
            queue.put(item)
        queue.put(sentinel)

    # start producer (in a background thread)
    import threading
    thread = threading.Thread(target=producer)
    thread.daemon = True

```

```

thread.start()

# run as consumer (read items from queue, in current thread)
item = queue.get()
while item is not sentinel:
    yield item
    queue.task_done()
    item = queue.get()

```

```

In [13]: from nolearn.lasagne import NeuralNet
        from lasagne.layers import DenseLayer, ReshapeLayer, Upscale2DLayer, Conv2DLayer, InputLayer, MaxPool2DLayer, get_all_params, batch_norm
        import numpy as np
        from lasagne.nonlinearities import softmax, leaky_rectify
        from lasagne.updates import nesterov_momentum
        from nolearn.lasagne import NeuralNet, BatchIterator, PrintLayerInfo, objective
        from nolearn.lasagne import TrainSplit
        from common import EarlyStopping, EndTrainingFromEarlyStopping
        from lasagne.objectives import categorical_crossentropy, aggregate
        import cPickle as pickle
        from sklearn import metrics
        import time, logging, logging.config, logging.handlers

```

Couldn't import dot\_parser, loading of dot files will not be possible.

Using gpu device 0: GeForce GTX 960 (CNMeM is disabled, CuDNN 4004)

In [14]: try:

```

        from lasagne.layers.dnn import Conv2DDNNLayer, MaxPool2DDNNLayer
        def conv_2_layer_stack(top, num_filters):
            conv1 = batch_norm(Conv2DDNNLayer(top, num_filters, (3, 3), stride=1, pad=1, nonlinearity=leaky_rectify))
            conv2 = batch_norm(Conv2DDNNLayer(conv1, num_filters, (3, 3), stride=1, pad=1, nonlinearity=leaky_rectify))
            return MaxPool2DDNNLayer(conv2, (2, 2), 2)

        def conv_4_layer_stack(top, num_filters):
            conv1 = batch_norm(Conv2DDNNLayer(top, num_filters, (3, 3), stride=1, pad=0, nonlinearity=leaky_rectify))
            conv2 = batch_norm(Conv2DDNNLayer(conv1, num_filters, (3, 3), stride=1, pad=0, nonlinearity=leaky_rectify))
            conv3 = batch_norm(Conv2DDNNLayer(conv2, num_filters, (3, 3), stride=1, pad=0, nonlinearity=leaky_rectify))
            conv4 = batch_norm(Conv2DDNNLayer(conv3, num_filters, (3, 3), stride=1, pad=0, nonlinearity=leaky_rectify))
            return MaxPool2DDNNLayer(conv4, (2, 2), 2)

        def conv_6_layer_stack(top, num_filters):
            conv1 = batch_norm(Conv2DDNNLayer(top, num_filters, (3, 3), stride=1, pad=1, nonlinearity=leaky_rectify))
            conv2 = batch_norm(Conv2DDNNLayer(conv1, num_filters, (3, 3), stride=1, pad=1, nonlinearity=leaky_rectify))
            conv3 = batch_norm(Conv2DDNNLayer(conv2, num_filters, (3, 3), stride=1, pad=1, nonlinearity=leaky_rectify))
            conv4 = batch_norm(Conv2DDNNLayer(conv3, num_filters, (3, 3), stride=1, pad=1, nonlinearity=leaky_rectify))
            conv5 = batch_norm(Conv2DDNNLayer(conv4, num_filters, (3, 3), stride=1, pad=1, nonlinearity=leaky_rectify))
            conv6 = batch_norm(Conv2DDNNLayer(conv5, num_filters, (3, 3), stride=1, pad=1, nonlinearity=leaky_rectify))
            return MaxPool2DLayer(conv6, (2, 2), 2)

    except ImportError:
        def conv_2_layer_stack(top, num_filters):
            conv1 = batch_norm(Conv2DLayer(top, num_filters, (3, 3), stride=1, pad=1, nonlinearity=leaky_rectify))
            conv2 = batch_norm(Conv2DLayer(conv1, num_filters, (3, 3), stride=1, pad=1, nonlinearity=leaky_rectify))
            return MaxPool2DLayer(conv2, (2, 2), 2)

```

```

def conv_4_layer_stack(top, num_filters):
    conv1 = batch_norm(Conv2DLayer(top, num_filters, (3, 3), stride=1, pad=0, nonlineari
    conv2 = batch_norm(Conv2DLayer(conv1, num_filters, (3, 3), stride=1, pad=0, nonlineari
    conv3 = batch_norm(Conv2DLayer(conv2, num_filters, (3, 3), stride=1, pad=0, nonlineari
    conv4 = batch_norm(Conv2DLayer(conv3, num_filters, (3, 3), stride=1, pad=0, nonlineari
    return MaxPool2DLayer(conv4, (2, 2), 2)

def conv_6_layer_stack(top, num_filters):
    conv1 = batch_norm(Conv2DLayer(top, num_filters, (3, 3), stride=1, pad=1, nonlineari
    conv2 = batch_norm(Conv2DLayer(conv1, num_filters, (3, 3), stride=1, pad=1, nonlineari
    conv3 = batch_norm(Conv2DLayer(conv2, num_filters, (3, 3), stride=1, pad=1, nonlineari
    conv4 = batch_norm(Conv2DLayer(conv3, num_filters, (3, 3), stride=1, pad=1, nonlineari
    conv5 = batch_norm(Conv2DLayer(conv4, num_filters, (3, 3), stride=1, pad=1, nonlineari
    conv6 = batch_norm(Conv2DLayer(conv5, num_filters, (3, 3), stride=1, pad=1, nonlineari
    return MaxPool2DLayer(conv6, (2, 2), 2)

In [15]: input_layer = InputLayer((None, 1, input_volume_shape[0], input_volume_shape[1]))
conv_stack_1 = conv_2_layer_stack(input_layer, 32)
conv_stack_2 = conv_2_layer_stack(conv_stack_1, 64)
conv_stack_3 = conv_4_layer_stack(conv_stack_2, 128)
conv_stack_4 = conv_4_layer_stack(conv_stack_3, 256)
dropout17 = DropoutLayer(conv_stack_4, p=0.5)
dense18 = DenseLayer(dropout17, 2048, nonlinearity=leaky_rectify)
dropout19 = DropoutLayer(dense18, p=0.5)
dense20 = DenseLayer(dropout19, 2048, nonlinearity=leaky_rectify)
softmax21 = DenseLayer(dense20, 10, nonlinearity=softmax)

```

## 1.1 Quality of Life Functions

```

In [16]: if not os.path.exists("logs"):
    os.mkdir("logs")
logging.config.fileConfig("logging-training.conf")

def regularization_objective(layers, lambda1=0., lambda2=0., *args, **kwargs):
    # default loss
    losses = objective(layers, *args, **kwargs)
    # get layer weights except for the biases
    weights = get_all_params(layers[-1], regularizable=True)
    regularization_term = 0.0
    # sum of abs weights for L1 regularization
    if lambda1 != 0.0:
        sum_abs_weights = sum([abs(w).sum() for w in weights])
        regularization_term += (lambda1 * sum_abs_weights)
    # sum of squares (sum(theta^2))
    if lambda2 != 0.0:
        sum_squared_weights = (1 / 2.0) * sum([(w ** 2).sum() for w in weights])
        regularization_term += (lambda2 * sum_squared_weights)
    # add weights to regular loss
    losses += regularization_term
    return losses

def eval_regularization(net):
    if net.objective_lambda1 == 0 and net.objective_lambda2 == 0:
        return 0

```

```

# check the loss if the regularization term is not overpowering the loss
weights = get_all_params(net.layers_[-1], regularizable=True)
# sum of abs weights for L1 regularization
sum_abs_weights = sum([abs(w).sum() for w in weights])
# sum of squares (sum(theta^2))
sum_squared_weights = (1 / 2.0) * sum([(w ** 2).sum() for w in weights])
# add weights to regular loss
regularization_term = (net.objective_lambda1 * sum_abs_weights) \
    + (net.objective_lambda2 * sum_squared_weights)
return regularization_term

def print_regularization_term(net):
    if net.objective_lambda1 > 0.0 or net.objective_lambda2 > 0.0:
        regularization_term = eval_regularization(net)
        print "Regularization term: {}".format(regularization_term.eval())

def validation_set_loss(_net, _X, _y):
    """We need this to track the validation loss"""
    _yb = _net.predict_proba(_X)
    _y_pred = np.argmax(_yb, axis=1)
    _acc = metrics.accuracy_score(_y, _y_pred)
    loss = aggregate(categorical_crossentropy(_yb, _y))
    loss += eval_regularization(_net)
    return loss, _acc

def store_model(model_file_name, net):
    directory_name = os.path.dirname(model_file_name)
    model_file_name = os.path.basename(model_file_name)
    if not os.path.exists(directory_name):
        os.makedirs(directory_name)
    # write model
    output_model_file_name = os.path.join(directory_name, model_file_name)
    start_write_time = time.time()
    if os.path.isfile(output_model_file_name):
        os.remove(output_model_file_name)
    with open(output_model_file_name, 'wb') as experiment_model:
        pickle.dump(net, experiment_model)
    total_write_time = time.time() - start_write_time
    m, s = divmod(total_write_time, 60)
    h, m = divmod(m, 60)
    logging.log(logging.INFO, "Duration of saving to disk: %0d:%02d:%02d", h, m, s)

def write_validation_loss_and_store_best(validation_file_name, best_weights_file_name,
                                         net, X_val, y_val, best_vloss, best_acc):
    # write validation loss
    start_validate_time = time.time()
    vLoss, vAcc = validation_set_loss(net, X_val, y_val)
    loss = vLoss.eval()
    current_epoch = net.train_history_[-1]['epoch']
    with open(validation_file_name, 'a') as validation_file:
        validation_file.write("{} {}, {} \n".format(current_epoch, loss, vAcc))

```

```

total_validate_time = time.time() - start_validate_time
m, s = divmod(total_validate_time, 60)
h, m = divmod(m, 60)
logging.log(logging.INFO, "Duration of validation: %0d:%02d:%02d", h, m, s)

# store best weights here
if loss < best_vloss:
    start_bw_time = time.time()
    best_vloss = loss
    best_acc = vAcc
    with open(best_weights_file_name, 'wb') as best_model_file:
        pickle.dump(net.get_all_params_values(), best_model_file, -1)

return best_vloss, best_acc

```

## 1.2 Define validation set

```

In [17]: val_dir = "/media/dylan/Science/Kaggle-Data/distracted_drivers/val/"
        X_val, y_val = image_gen_from_dir(val_dir, 40, 10, size=input_volume_shape).next()
        X_val = X_val.reshape(-1, 1, input_volume_shape[0], input_volume_shape[1])

```

## 1.3 CNN

```

In [18]: lambda1 = 0.0
        lambda2 = 5e-4

net = NeuralNet(
    layers=softmax21,
    max_epochs=1,
    update=nesterov_momentum,
    update_learning_rate=0.0001,
    update_momentum = 0.9,
    # update=adam,
    on_epoch_finished=[
        EarlyStopping(patience=2000)
    ],
    on_training_finished=[
        EndTrainingFromEarlyStopping()
    ],
    objective=regularization_objective,
    objective_lambda2=lambda2,
    objective_lambda1=lambda1,
    batch_iterator_train=BatchIterator(batch_size=100),
    train_split=TrainSplit(
        eval_size=0.25),
    # train_split=TrainSplit(eval_size=0.0),
    verbose=3,
)

In [ ]: p = PrintLayerInfo()
        net.initialize()
        # p(net)

In [ ]: image_gen = image_gen_from_dir(data_dir, 10, 10, size=input_volume_shape)
        gen = random_aug_gen(image_gen, random_aug)

```

```

threaded_gen = threaded_generator(gen, num_cached=100)

ops_every = 100
dir_name = 'net.vgg.large.12.5e4'
validation_file_name = "{} / vloss-{}.txt".format(dir_name, dir_name)
model_file_name = "{} / {}.pickle".format(dir_name, dir_name)
best_weights_file_name = "{} / bw-{}.weights".format(dir_name, dir_name)
best_acc = 0.0
best_vloss = np.inf

start_time = time.time()
try:
    for step, (inputs, targets) in enumerate(threaded_gen):
        shape = inputs.shape
        net.fit(inputs.reshape(shape[0],1, shape[1], shape[2]), targets)
        if (step + 1) % ops_every == 0:
            print_regularization_term(net)
            store_model(model_file_name, net)
            # center validation
            best_vloss, best_acc = write_validation_loss_and_store_best(
                validation_file_name, best_weights_file_name, net, X_val, y_val, best_vloss, be

except StopIteration:
    # terminate if already early stopping
    with open("net.vgg.large.pickle", 'wb') as writer:
        pickle.dump(net, writer)
    total_time = time.time() - start_time
    print("Training successful by early stopping. Elapsed: {}".format(total_time))

# Neural Network with 8964778 learnable parameters

## Layer information

```

name	size	total	cap.Y	cap.X	cov.Y	cov.X	filter Y	filter X	f
-----	-----	-----	-----	-----	-----	-----	-----	-----	-----
InputLayer	1x128x128	16384	100.00	100.00	100.00	100.00	128	128	
Conv2DDNNLayer	32x128x128	524288	100.00	100.00	2.34	2.34	3	3	
BatchNormLayer	32x128x128	524288	100.00	100.00	100.00	100.00	128	128	
NonlinearityLayer	32x128x128	524288	100.00	100.00	100.00	100.00	128	128	
Conv2DDNNLayer	32x128x128	524288	100.00	100.00	100.00	100.00	128	128	
BatchNormLayer	32x128x128	524288	100.00	100.00	100.00	100.00	128	128	
NonlinearityLayer	32x128x128	524288	100.00	100.00	100.00	100.00	128	128	
MaxPool2DDNNLayer	32x64x64	131072	100.00	100.00	100.00	100.00	128	128	
Conv2DDNNLayer	64x64x64	262144	100.00	100.00	100.00	100.00	128	128	
BatchNormLayer	64x64x64	262144	100.00	100.00	100.00	100.00	128	128	
NonlinearityLayer	64x64x64	262144	100.00	100.00	100.00	100.00	128	128	
Conv2DDNNLayer	64x64x64	262144	100.00	100.00	100.00	100.00	128	128	
BatchNormLayer	64x64x64	262144	100.00	100.00	100.00	100.00	128	128	
NonlinearityLayer	64x64x64	262144	100.00	100.00	100.00	100.00	128	128	
MaxPool2DDNNLayer	64x32x32	65536	100.00	100.00	100.00	100.00	128	128	
Conv2DDNNLayer	128x30x30	115200	100.00	100.00	100.00	100.00	128	128	
BatchNormLayer	128x30x30	115200	100.00	100.00	100.00	100.00	128	128	
NonlinearityLayer	128x30x30	115200	100.00	100.00	100.00	100.00	128	128	
Conv2DDNNLayer	128x28x28	100352	100.00	100.00	100.00	100.00	128	128	



BatchNormLayer	128x28x28	100352	100.00	100.00	100.00	100.00	128	128
NonlinearityLayer	128x28x28	100352	100.00	100.00	100.00	100.00	128	128
Conv2DDNNLayer	128x26x26	86528	100.00	100.00	100.00	100.00	128	128
BatchNormLayer	128x26x26	86528	100.00	100.00	100.00	100.00	128	128
NonlinearityLayer	128x26x26	86528	100.00	100.00	100.00	100.00	128	128
Conv2DDNNLayer	128x24x24	73728	100.00	100.00	100.00	100.00	128	128
BatchNormLayer	128x24x24	73728	100.00	100.00	100.00	100.00	128	128
NonlinearityLayer	128x24x24	73728	100.00	100.00	100.00	100.00	128	128
MaxPool2DDNNLayer	128x12x12	18432	100.00	100.00	100.00	100.00	128	128
Conv2DDNNLayer	256x10x10	25600	100.00	100.00	100.00	100.00	128	128
BatchNormLayer	256x10x10	25600	100.00	100.00	100.00	100.00	128	128
NonlinearityLayer	256x10x10	25600	100.00	100.00	100.00	100.00	128	128
Conv2DDNNLayer	256x8x8	16384	100.00	100.00	100.00	100.00	128	128
BatchNormLayer	256x8x8	16384	100.00	100.00	100.00	100.00	128	128
NonlinearityLayer	256x8x8	16384	100.00	100.00	100.00	100.00	128	128
Conv2DDNNLayer	256x6x6	9216	100.00	100.00	100.00	100.00	128	128
BatchNormLayer	256x6x6	9216	100.00	100.00	100.00	100.00	128	128
NonlinearityLayer	256x6x6	9216	100.00	100.00	100.00	100.00	128	128
Conv2DDNNLayer	256x4x4	4096	100.00	100.00	100.00	100.00	128	128
BatchNormLayer	256x4x4	4096	100.00	100.00	100.00	100.00	128	128
NonlinearityLayer	256x4x4	4096	100.00	100.00	100.00	100.00	128	128
MaxPool2DDNNLayer	256x2x2	1024	100.00	100.00	100.00	100.00	128	128
DropoutLayer	256x2x2	1024	100.00	100.00	100.00	100.00	128	128
DenseLayer	2048	2048	100.00	100.00	100.00	100.00	128	128
DropoutLayer	2048	2048	100.00	100.00	100.00	100.00	128	128
DenseLayer	2048	2048	100.00	100.00	100.00	100.00	128	128
DenseLayer	10	10	100.00	100.00	100.00	100.00	128	128

#### Explanation

X, Y: image dimensions  
 cap.: learning capacity  
 cov.: coverage of image  
 magenta: capacity too low (<1/6)  
 cyan: image coverage too high (>100%)  
 red: capacity too low and coverage too high

epoch	train loss	valid loss	train/val	valid acc	dur
1	3.23474	2.47064	1.30927	0.06667	0.92s
2	3.13003	2.47110	1.26665	0.03333	0.90s
3	2.99001	2.46972	1.21066	0.06667	1.01s
4	3.25569	2.46959	1.31831	0.10000	0.88s
5	3.13186	2.47047	1.26772	0.06667	0.90s
6	3.10504	2.46985	1.25718	0.10000	0.91s
7	3.49868	2.48008	1.41071	0.06667	0.90s
8	3.24035	2.47678	1.30829	0.10000	0.90s
9	3.00858	2.48594	1.21024	0.10000	0.95s
10	3.07501	2.48129	1.23928	0.06667	0.90s
11	3.01663	2.47993	1.21642	0.13333	0.92s
12	3.03845	2.45505	1.23763	0.06667	0.92s
13	2.91237	2.50508	1.16259	0.10000	0.91s
14	2.85795	2.46097	1.16131	0.16667	0.90s
15	2.80923	2.47287	1.13602	0.13333	0.91s

16	3.01027	2.48820	1.20982	0.10000	0.90s
17	2.80844	2.50207	1.12245	0.10000	0.90s
18	2.86454	2.44450	1.17183	0.10000	0.89s
19	2.79950	2.51300	1.11401	0.03333	0.90s
20	3.02547	2.56349	1.18022	0.03333	0.91s
21	2.82010	2.53362	1.11307	0.10000	0.89s
22	2.82993	2.42642	1.16630	0.16667	0.89s
23	3.12718	2.56695	1.21825	0.10000	0.90s
24	2.77314	2.54414	1.09001	0.10000	0.91s
25	2.95101	2.60550	1.13261	0.10000	0.90s
26	2.85654	2.49399	1.14537	0.10000	0.97s
27	2.84319	2.41811	1.17579	0.16667	0.95s
28	2.72599	2.58565	1.05428	0.10000	0.89s
29	3.06946	2.59261	1.18393	0.10000	0.90s
30	2.81219	2.51542	1.11798	0.10000	0.91s
31	3.06138	2.54847	1.20126	0.13333	0.91s
32	2.82345	2.49851	1.13006	0.16667	0.90s
33	2.90232	2.50254	1.15975	0.10000	0.89s
34	2.75007	2.53666	1.08413	0.06667	0.90s
35	2.95241	2.62377	1.12525	0.06667	0.91s
36	3.04474	2.60348	1.16949	0.00000	0.92s
37	2.92750	2.50788	1.16732	0.10000	0.90s
38	2.53108	2.58230	0.98017	0.13333	0.90s
39	3.09970	2.52990	1.22523	0.10000	0.90s
40	2.67083	2.54307	1.05024	0.10000	0.90s
41	2.88905	2.44300	1.18258	0.13333	0.92s
42	3.00076	2.63735	1.13779	0.03333	0.91s
43	2.84186	2.66385	1.06683	0.10000	0.91s
44	2.83282	2.52885	1.12020	0.03333	0.90s
45	2.93888	2.58979	1.13479	0.13333	0.91s
46	2.75568	2.49318	1.10529	0.20000	0.89s
47	2.78819	2.60775	1.06919	0.10000	0.90s
48	3.00252	2.60190	1.15397	0.00000	0.89s
49	2.85724	2.55592	1.11789	0.13333	0.89s
50	2.94593	2.53239	1.16330	0.10000	0.89s
51	3.03594	2.52270	1.20345	0.10000	0.90s
52	2.81859	2.55127	1.10478	0.16667	0.89s
53	2.77743	2.54951	1.08940	0.10000	0.91s
54	2.94093	2.44697	1.20186	0.23333	0.91s
55	2.95946	2.48883	1.18910	0.03333	0.91s
56	2.76457	2.49127	1.10970	0.06667	0.91s
57	2.81538	2.52222	1.11623	0.06667	0.90s
58	2.63187	2.62455	1.00279	0.10000	0.90s
59	2.86669	2.62714	1.09118	0.10000	0.90s
60	2.80323	2.58619	1.08392	0.16667	0.91s
61	2.78928	2.59084	1.07659	0.10000	0.90s
62	2.78297	2.49973	1.11331	0.10000	0.90s
63	2.99439	2.51418	1.19100	0.10000	0.91s
64	2.85013	2.54657	1.11920	0.10000	0.90s
65	2.75358	2.55932	1.07590	0.13333	0.90s
66	2.81624	2.56078	1.09976	0.10000	0.91s
67	2.81874	2.50096	1.12706	0.10000	0.91s
68	2.84398	2.63137	1.08080	0.20000	0.91s
69	2.84880	2.51919	1.13084	0.06667	0.91s

70	2.93559	2.53103	1.15984	0.06667	0.90s
71	2.70180	2.45847	1.09898	0.06667	0.90s

In [ ]: